

Deep Learning

Assignment 5

Word Embedding

Aim

To develop a neural network model to perform classification on the Reuters dataset with both trainable and loaded embedding layer and to test the accuracy of each model. Finally LSTM layers are added to the model with good accuracy to make it a further better model.

Problem Statement

The IMDB dataset consists of the reviews of movies. These reviews are to be classified based on the terms used in the reviews by a neural network model.

Proposed Solution

Dataset

The Reuters dataset is used the most in the past decades. It consists of 21578 topics of texts without its topics and typographical errors. The purpose of this model is to classify the data according to its topics. These texts fall under 46 such categories that are to be classified.

Embedding layer

An embedding layer is created to which the dataset is passed. The dimensions of the embedding layer matrices are specified. This layer creates a unique 2D vector for every word in the given input. By this way it makes the process of language processing easier.

Dense Layer

Finally dense layers are attached to the model with ReLU activation. Since the output of the embedding layer has to be passed to the dense layer, the outputs of the embedding layer are flattened. Finally the model is built with the optimizer RMSProp and binary cross entropy loss and compiled.

Data preprocessing

The data was preprocessed by splitting it into test, train and validation datasets. Then the labels were converted to one-hot encodings. Since the input

datasets could be of various string lengths they were padded with zeroes for a definite length.

Glove embedding

The Glove embedding layer was used for the loaded embedding. GloVe stands for global vectors for word representation. It is an unsupervised learning algorithm for generating word embedding by aggregating global word-word co-occurrence matrix from a corpus. The resulting embedding show interesting linear substructures of the word in vector space.

LSTM layer

It stands for Long Short Term Memory. It consists of feedback connections that help it to remember the sequences of input. Thus it can give a better accuracy than other deep learning models.

Evaluation

Trainable model

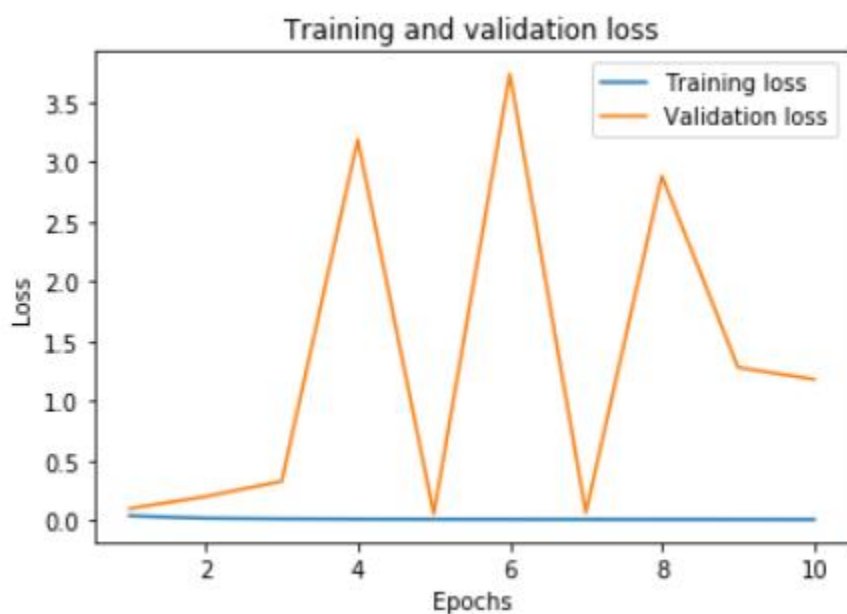
Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 2000, 8)	80000
batch_normalization_1 (Batch Normalization)	(None, 2000, 8)	32
flatten_1 (Flatten)	(None, 16000)	0
dense_1 (Dense)	(None, 512)	8192512
dense_2 (Dense)	(None, 46)	23598
Total params: 8,296,142		
Trainable params: 8,296,126		
Non-trainable params: 16		

Fitting the model

```
Epoch 6/10
8083/8083 [=====] - 23s 3ms/step - loss: 0.0088 -
acc: 0.9978 - val_loss: 3.7361 - val_acc: 0.9783
Epoch 7/10
8083/8083 [=====] - 22s 3ms/step - loss: 0.0083 -
acc: 0.9979 - val_loss: 0.0638 - val_acc: 0.9897
Epoch 8/10
```

```
8083/8083 [=====] - 24s 3ms/step - loss: 0.0080 -  
acc: 0.9980 - val_loss: 2.8806 - val_acc: 0.9783  
Epoch 9/10  
8083/8083 [=====] - 23s 3ms/step - loss: 0.0076 -  
acc: 0.9980 - val_loss: 1.2796 - val_acc: 0.9785  
Epoch 10/10  
8083/8083 [=====] - 20s 3ms/step - loss: 0.0071 -  
acc: 0.9981 - val_loss: 1.1797 - val_acc: 0.9783
```



Test accuracy – 97%

```
2246/2246 [=====] - 1s 284us/step  
[1.155681626985674, 0.9783766269683838]
```


Glove model

Model: "sequential_3"

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, 2000, 100)	1000000
flatten_3 (Flatten)	(None, 200000)	0
dense_5 (Dense)	(None, 512)	102400512
dropout_3 (Dropout)	(None, 512)	0
dense_6 (Dense)	(None, 46)	23598
Total params: 103,424,110		
Trainable params: 103,424,110		
Non-trainable params: 0		

Fitting the model

```
Epoch 6/10  
8083/8083 [=====] - 178s 22ms/step - loss: 0.2519  
- acc: 0.9783 - val_loss: 0.2302 - val_acc: 0.9783  
Epoch 7/10  
8083/8083 [=====] - 179s 22ms/step - loss: 0.2109  
- acc: 0.9783 - val_loss: 0.1938 - val_acc: 0.9783  
Epoch 8/10  
8083/8083 [=====] - 170s 21ms/step - loss: 0.1783  
- acc: 0.9783 - val_loss: 0.1649 - val_acc: 0.9783  
Epoch 9/10  
8083/8083 [=====] - 172s 21ms/step - loss: 0.1525  
- acc: 0.9783 - val_loss: 0.1423 - val_acc: 0.9783  
Epoch 10/10  
8083/8083 [=====] - 168s 21ms/step - loss: 0.1325  
- acc: 0.9783 - val_loss: 0.1247 - val_acc: 0.9783
```



Test accuracy – 97%

```
2246/2246 [=====] - 3s 1ms/step
[0.12384231522586446, 0.9782605171203613]
```

LSTM

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, None, 8)	80000
lstm_1 (LSTM)	(None, 128)	70144
dense_1 (Dense)	(None, 46)	5934
Total params: 156,078		
Trainable params: 156,078		
Non-trainable params: 0		

Fitting the model

Train on 8083 samples, validate on 899 samples

Epoch 1/5


```
8083/8083 [=====] - 475s 59ms/step - loss: 0.0862
- acc: 0.9755 - val_loss: 0.0741 - val_acc: 0.9783
```

Epoch 2/5

```

8083/8083 [=====] - 302s 37ms/step - loss: 0.0723
- acc: 0.9783 - val_loss: 0.0739 - val_acc: 0.9783
Epoch 3/5
8083/8083 [=====] - 315s 39ms/step - loss: 0.0721
- acc: 0.9783 - val_loss: 0.0741 - val_acc: 0.9783
Epoch 4/5
8083/8083 [=====] - 311s 38ms/step - loss: 0.0720
- acc: 0.9783 - val_loss: 0.0736 - val_acc: 0.9783
Epoch 5/5
8083/8083 [=====] - 313s 39ms/step - loss: 0.0719
- acc: 0.9783 - val_loss: 0.0740 - val_acc: 0.9783

```



Testing accuracy – 97%

```

2246/2246 [=====] - 17s 8ms/step
[0.07250244192653008, 0.9782605171203613]

```

2 LSTM

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
embedding_2 (Embedding)	(None, None, 8)	80000
=====		
lstm_2 (LSTM)	(None, None, 25)	3400
=====		
lstm_3 (LSTM)	(None, 25)	5100
=====		
dense_2 (Dense)	(None, 46)	1196
=====		

Total params: 89,696
Trainable params: 89,696
Non-trainable params: 0

Fitting the model

Train on 8083 samples, validate on 899 samples

Epoch 1/5

8083/8083 [=====] - 285s 35ms/step - loss: 0.1623
- acc: 0.9573 - val_loss: 0.0737 - val_acc: 0.9783

Epoch 2/5

8083/8083 [=====] - 281s 35ms/step - loss: 0.0719
- acc: 0.9783 - val_loss: 0.0735 - val_acc: 0.9783

Epoch 3/5


8083/8083 [=====] - 282s 35ms/step - loss: 0.0733
- acc: 0.9783 - val_loss: 0.0735 - val_acc: 0.9783

Epoch 4/5

8083/8083 [=====] - 285s 35ms/step - loss: 0.0718
- acc: 0.9783 - val_loss: 0.0735 - val_acc: 0.9783

Epoch 5/5

8083/8083 [=====] - 284s 35ms/step - loss: 0.0718
- acc: 0.9783 - val_loss: 0.0736 - val_acc: 0.9783



Testing accuracy – 97%

2246/2246 [=====] - 10s 4ms/step
[0.07206905480616876, 0.9782605171203613]

Conclusion

All the models seemed to give the same high accuracy. This could be because the datasets were large and the models got trained well.

References

Code given by the Professor.

Keras code for reuters classification